

UNIVERSITY of CALIFORNIA
SANTA CRUZ

**SOLVING BOUNDARY VALUE PROBLEMS EFFICIENTLY
USING COMPUTERS [working title]**

A thesis submitted in partial satisfaction
of the requirements for the degree of

BACHELOR OF SCIENCE

in

PHYSICS

by

Daniel A. Bittman
May 11, 2016

The thesis of Daniel A. Bittman is approved by:

Professor Jason Nielsen
Thesis Advisor

Professor David P. Belanger
Thesis Coordinator

Professor Robert Johnson
Chair, Department of Physics

COPYRIGHT AND SOFTWARE LICENSE

THIS DOCUMENT AND SOFTWARE DESCRIBED HEREIN ARE

COPYRIGHT © 2016
DANIEL A. BITTMAN

THE SOFTWARE DESCRIBED HEREIN, INCLUDING EXCERPTS AND OTHER CODE, IS
PROVIDED WITH THE FOLLOWING LICENSE:

MIT License

Copyright © 2016 Daniel A. Bittman

Permission is hereby granted, free of charge, to any person obtaining a copy of this software and associated documentation files (the "Software"), to deal in the Software without restriction, including without limitation the rights to use, copy, modify, merge, publish, distribute, sublicense, and/or sell copies of the Software, and to permit persons to whom the Software is furnished to do so, subject to the following conditions:

The above copyright notice and this permission notice shall be included in all copies or substantial portions of the Software.

THE SOFTWARE IS PROVIDED "AS IS", WITHOUT WARRANTY OF ANY KIND, EXPRESS OR IMPLIED, INCLUDING BUT NOT LIMITED TO THE WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE AND NONINFRINGEMENT. IN NO EVENT SHALL THE AUTHORS OR COPYRIGHT HOLDERS BE LIABLE FOR ANY CLAIM, DAMAGES OR OTHER LIABILITY, WHETHER IN AN ACTION OF CONTRACT, TORT OR OTHERWISE, ARISING FROM, OUT OF OR IN CONNECTION WITH THE SOFTWARE OR THE USE OR OTHER DEALINGS IN THE SOFTWARE.

DEDICATION (TODO)

ACKNOWLEDGEMENTS (TODO)

Contents

1	Introduction	2
1.1	Discretizing Poisson's Equation	3
1.2	Performance Considerations	4
2	Methods	6
2.1	Jacobi Iteration	6
2.2	Successive Over-Relaxation	7
2.3	Design of Solver	8
2.4	Configurations and Mapping to Electrostatics	8
2.5	Verifying Correctness	9
2.6	Potential From Electric Dipole	11
2.7	Optimization and Multithreading	12
2.8	SIMD Instructions	14
2.9	Multithreading	15
2.10	Benchmarking Performance	15
3	Results	17
3.1	The $\sin(x)$ Potential	17
3.2	Fitting to a $1/r^2$ Dipole Potential	19
3.3	Performance	22
4	Conclusion	27
4.1	Correctness	27

4.2	Performance	27
4.3	Future Work	28
	Appendices	30
	A Program Usage	30
	B Computer Program Optimization	31
B.1	Lazy Programs Are Better	31
B.2	Freebies	32
B.3	A Model of Processor Design	34
B.3.1	Caching	36
B.4	Structure of Arrays	37
B.5	Multithreading	37
B.6	Vectorization	37
B.6.1	Branch Reduction	37
	C Program Implementation	38

List of Figures

1	Visualization of grid, with indicies.	3
2	Result of simulation for the $\sin(x)$ potential.	18
3	Analytical result for the $\sin(x)$ potential.	19
4	Difference between the calculated and analytical results for the $\sin(x)$ potential.	20
5	Result of the simulation for the dipole configuration, including contour lines.	21
6	Fitting a $1/r^2$ curve to the calculated dipole potential.	22
7	Result of program outputting the electric field for the dipole configuration.	23
8	Different grid sizes for the $\sin(x)$ configuration, and their performance characteristics.	24
9	Performance results for a varying number of threads in non-SIMD mode. These tests were done on computer B, using the $\sin(x)$ potential with a grid size of 300 by 300.	25
10	Error results for a varying number of threads in non-SIMD mode. These tests were done on computer B, using the $\sin(x)$ potential with a grid size of 300 by 300.	26
11	View of programming that programmers like to have (a), versus the the more realistic and complicated view (b).	35

List of Tables

1	Unoptimized and optimized assembly from the sum-a-lot-of-numbers example.	34
---	---	----

Abstract

There are many problems in physics which involve solving a set of equations which define behavior inside some region. One such set of problems, among many, is solving electrostatics problems inside a region given some boundary conditions by utilizing Poisson's equation. Such problems, and other problems, are commonly found in simple and analytically solvable forms in undergraduate physics classes. Unfortunately, these problems can become intractable to solve analytically very quickly as their complexity increases. Complex boundary conditions in strangely shaped regions can easily become impossibly difficult to solve. The solution to this problem can be to model it on a computer in order to compute a numerical solution in place of an analytical one.

I have written a software package which is both a C library which can solve such a given problem and a frontend Python program which can parse a human-readable configuration that describes such a boundary value problem and set up the correct initial state such that the C library can then solve it, allowing the Python program to display the final result as well as calculate statistics about the solving process. I have utilized effective programming techniques designed for current generation hardware 64-bit x86 processors in order to improve the performance of the program, in addition to implementing a multi-threaded version of the solving program to explore the use of multicore processors in solving these problems.

I found that the program correctly matched an analytical solution to a simple test problem, and that it correctly predicted the $1/r^2$ potential found from an electric dipole. I was able to optimize the program so that it efficiently found these results, and I was able to successfully leverage multithreading to significantly increase the performance, albeit with a degradation in simulation accuracy.

1 Introduction

Boundary value problems, problems defined by a partial differential equation and a set of constraints on the solution, are found in many areas of physics. These include problems such as finding the temperature of a rod that has a given temperature on either end, determining the electric potential inside a region given the potential around the boundary, or determining the flow of an incompressible fluid. As the complexity of these problems increases, however, the feasibility of solving them exactly by hand quickly decreases. Instead, we have the option to solve them with iterative methods using computers. These types of problems are well-suited for computers because they are easily parallelizable [8]. I have developed an implementation of such a solver for electrostatics and fluid flow utilizing the parallelizability of the problem and designed it to be efficient, taking advantage of the design of modern CPUs.

Although many differential equations are typically used in boundary value problems, I will limit the focus to the Poisson equation. Named after Simon Denis Poisson, the Poisson equation is used in a number of typical boundary value problems, including electrostatics. In 2-dimensional Euclidean space it is[1],

$$\nabla^2\phi = f,$$

where f is a real-valued function inside the space and ϕ is the steady-state solution for the potential in the space. However, since this solution will be calculated numerically on a computer, we must discretize the equation so that it can be used to represent a solution on a 2-dimensional grid.

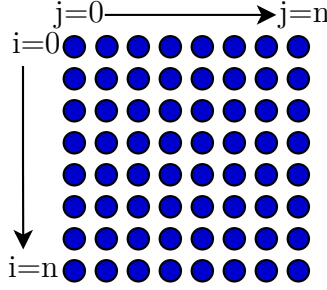


Figure 1: Visualization of grid, with indicies.

1.1 Discretizing Poisson's Equation

Consider a grid of size n by n as in figure 1¹, with values represented by $u_{i,j}$. We want to apply the Poisson equation to this u as follows:

$$\nabla^2 u = \frac{\partial^2 u}{\partial x^2} + \frac{\partial^2 u}{\partial y^2} = f,$$

however, u is not continuous, so the partial derivative for x becomes

$$\frac{\partial^2 u}{\partial x^2} \approx \frac{u_{i-1,j} - 2u_{i,j} + u_{i+1,j}}{h^2},$$

at a grid point (i,j) , where $h = 1/n$ is the width and height of a grid cell. The partial derivative for y is similar, using j instead of i . This gives the full discrete Poisson equation as[2][9],

$$\frac{u_{i-1,j} - 2u_{i,j} + u_{i+1,j}}{h^2} + \frac{u_{i,j-1} - 2u_{i,j} + u_{i,j+1}}{h^2} = f_{i,j}.$$

Since $u_{i,j}$ is what we are interested in, this is more useful after some rearranging:

$$h^2 f_{i,j} = -4u_{i,j} + u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1},$$

¹In general, it is not difficult to allow a rectangular grid, but for simplicity we will require the grid to be square.

$$u_{i,j} = \frac{1}{4}(u_{i-1,j} + u_{i+1,j} + u_{i,j-1} + u_{i,j+1} - h^2 f_{i,j}). \quad (1)$$

It can be seen from the top equation that this is now a linear algebra problem, requiring a method of solving n^2 dependent linear equations. This means that we can apply one of several standard iterative methods for solving the problem. This is an excellent application of using a computer to solve numerically a problem which is difficult to solve analytically.

1.2 Performance Considerations

While operations on a grid of values are parallelizable, there is a limit to the effectiveness of multi-threading the algorithm. Take, for example, an average modern processor with 4 cores each with a 256KB cache and a shared cache of 8MB. Each core can only operate on a grid of 256 by 256 32-bit values before exhausting its cache. A performance drop will occur if the grid size per core exceeds the core's cache size. For example, should the grid size exceed 1440 by 1440 32-bit values, the shared 8MB cache will be filled, resulting in a massive performance penalty. For this reason, naïvely splitting the grid into 4 areas and multi-threading the iteration may result in sub-par performance.

There are many other aspects of processor design that come into play when writing a program such as this. Multithreaded synchronization aside, there are many details of caching that need to be considered. Of course, optimizing a program is just like any experiment in that it is absolutely vital to apply normal scientific practices when doing such work. Many make the mistake of simply writing what they believe to be a more optimized version of a program without measuring the resulting performance and comparing it to the previous version.

I would expect to see the performance enhanced by the use of multithreading, but insignificantly due to the large amount of memory accesses made by the program throughout solving a given problem.

2 Methods

Since we are using a computer, an iterative method is feasible to be used as a numerical solution method. In this program, I have implemented two methods for solving these problems: Jacobi iteration, and successive over-relaxation (SOR).

2.1 Jacobi Iteration

This method is perhaps the simplest method of doing an iterative solution to this linear algebra problem. The procedure involves treating each equation as independent, and solving iteratively. This turns Equation 1 into[2]

$$u_{i,j}^{k+1} = \frac{1}{4}(u_{i-1,j}^k + u_{i+1,j}^k + u_{i,j-1}^k + u_{i,j+1}^k - h^2 f_{i,j}), \quad (2)$$

where k is the iteration number. This is saying that to calculate u everywhere, we must start with a guess for u , and then iteratively solve each element in u , which depends on its four nearest neighbors. One immediate unfortunate reality with Jacobi iteration presents itself: we must store a secondary copy of the entire grid because each calculation of $u_{i,j}^{k+1}$ strictly requires values from the previous iteration. Since we iterate over the grid in order to solve it, we cannot write the new value for the $k + 1$ iteration into that element of u , since we would be overwriting a value that is needed later. This unfortunately but necessarily increases the space required by a factor of n^2 .

The above equation does not yet include any consideration of boundary conditions. In fact, this is only valid for a region that is unbounded. The primary functionality of supporting solving statics problems requires support for dealing with Dirichlet boundary conditions. Dirichlet boundary conditions are defined such that the solution ϕ of a partial differential

equations in a region Ω has a value $\phi(\mathbf{r}) = y(\mathbf{r}) \forall \mathbf{r} \in \partial\Omega$. For this case, this means that if a Dirichlet boundary value condition is specified for a particular grid cell (x, y) , then its value is simply fixed to the value of the Dirichlet condition, and that cell does not need to be recalculated on successive iterations.

Another consideration is that the grid is of finite size. This means that when calculating $u_{i=0,j}$, for example, we cannot incorporate $u_{i-1,j}$ as Equation 2 would dictate. This is because that would try to get the -1 st row of the matrix u , which does not exist. A common method of dealing with this problem is “wrap around ” to the other side – essentially considering the region Ω to be the surface of a sphere, or to consider the space to be periodic. Because of the nature of the problems that I will be considering, I have chosen to consider an implicit Dirichlet boundary condition of zero everywhere outside of Ω , thus effectively limiting the calculation to be inside the region while clamping any values “outside” of Ω to zero. This is easy to implement and visualize as well: any value of $u_{i,j}$ for $i \geq n$, $i < 0$, $j \geq n$, or $j < 0$ is simply zero, always.

2.2 Successive Over-Relaxation

An alternate approach to iteratively solving this problem numerically is called successive over-relaxation (or SOR), and it has advantages over simple Jacobi iteration. It works by reusing previous calculations for the next iteration when calculating a given cell. It transforms Equation 2 into,

$$u_{i,j}^{k+1} = (1 - \omega) + \frac{\omega}{4} \left(u_{i-1,j}^{k+1} + u_{i+1,j}^k + u_{i,j-1}^{k+1} + u_{i,j+1}^k - h^2 f_{i,j} \right), \quad (3)$$

where ω is a tunable parameter called the *over-relaxtion parameter*. For a square lattice, a value of

$$\omega = \frac{2}{1 + \frac{\pi}{n}}$$

has the best rate of convergence[2]. There are two terms on the right-hand-side of equation 3 which refer to the $k + 1$ iteration. What this is doing is using the new values for the cells that have already been updated on this iteration in order to speed up the convergence.

2.3 Design of Solver

The solving program is written in C in order to leverage the full performance of modern computers and is implemented as a shared library (on UNIX) such that it can be linked with another program which sets up the initial conditions before calling the solving function. The C code that does the solving is hereby referred to as the engine, while any code that interfaces with the engine and provides a usable user interface is referred to as the frontend.

I have written an example frontend in Python, which is capable to parsing a configuration file which describes the values of f , the initial boundary conditions, the grid size, and a number of other details about a given desired setup (hereby collectively referred to as a configuration). I have written several example configurations which I have used to confirm the correctness of the solving program. For a more detailed technical description of the implementation, see Appendix C.

2.4 Configurations and Mapping to Electrostatics

These configuration files are designed to be easy to write. An example configuration file is,

```

gridsize 300
dirichlet 1,1 1,grid.len = math.sin(y * math.pi / grid.len)
cell 10,10 initial -100

```

The first line sets up a grid of size 300 by 300. The second line specifies that a Dirichlet boundary condition exists from cell (1,1) to (1,300) which takes on a value given by the expression at the end of that line. Any valid python expression may be put here, allowing access to complex mathematical expressions. Additionally, the start and end points do not have to define a line on a single axis, the program will do a linear interpolation to create a straight line from the start point to the end point. The third line sets the initial value of the cell (10,10) to a value of -100 . This is the command used to define $f_{i,j}$ in equation 1. An arbitrary number of these statements may be made inside a configuration file, allowing one to specify the value of f everywhere if need be, or create Dirichlet boundary conditions of arbitrary shapes.

The available commands for a configuration file map to various initial and boundary conditions for an electrostatics problem in a straight forward way: the command to specify a Dirichlet boundary condition is like creating a boundary of fixed potential. The command to specify the initial value of cells is like specifying an initial charge distribution inside the region.

2.5 Verifying Correctness

The clearest way to verify that the solver is working correctly is to test it out on a number of different configurations for which the correct solution is known. My first choice here was a simple electrostatics problem straight out of a textbook[3]:

Consider a square region with the bottom left corner at the origin with each side of length L . Let the potential along each side of the region be fixed at 0, except for

the bottom side, which has a potential $V(x) = V_0 \sin(x\pi/L)$. What is the potential V inside the region?

This problem describes solving Laplace's equation (which is Poisson's equation with $f = 0$), given some initial boundary conditions. It can be solved analytically using a standard separation of variables trick[1]. Assume

$$\phi = X(x)Y(y),$$

then

$$X''(x)/X(x) = -Y''(y)/Y(y) = -\lambda,$$

and by applying the boundary conditions, one comes to the solution[3]

$$V(x, y) = \frac{\sinh(y\pi/L)}{\sinh(\pi/L)} \sin(x\pi/L).$$

In order to compare the result of my program with the analytical solution, I added a feature to the Python frontend to accept a pre-made file containing a precalculated grid of the correct output by using the analytic solution. The frontend then reads this file and compares it to the result of the solver.

Since there is now a problem which has an analytical solution, I was able to use that to describe a metric for correctness of the simulation. I wrote a Python script that generated a 2-D grid of values based on the analytical solution above and had it write this data structure out to disk. When the solver program was given this "correct" file, it would then perform a root-mean-square error calculation,

$$e_{RMS} = \sqrt{\left(\frac{\sum_{i=0}^n \sum_{j=0}^n (u_{i,j} - c_{i,j})^2}{n^2} \right)}$$

where n is the length of one side of the grid, $u_{i,j}$ is the value of the calculated result at cell (i,j) , and $c_{i,j}$ is the value of the analytical result at cell (i,j) .

2.6 Potential From Electric Dipole

Another correctness check is a simulation of an electric dipole, since the potential from an electric dipole is well understood. If two equal and opposite charges are separated on an axis by distance d , then at a point p far away, the potential is,

$$V = C \left(\frac{1}{r_0} - \frac{1}{r_1} \right),$$

where r_0 and r_1 are the distances from each charge, and C is a constant. With some rearranging,

$$V = C \left(\frac{r_0 - r_1}{r_0 r_1} \right) = \frac{Cd \cos(\theta)}{r^2},$$

where r is the distance from the center of charge to the point, and θ is the angle that the line from the center of charge forms with the axis that the dipole is on[3]. The potential from an electric dipole thus falls off as $1/r^2$, which should be reflected in the solution given by the program if it is given a configuration describing a dipole.

If the solving program is correct, then if two charges are placed near each other, the numerically calculated potential should display this $1/r^2$ behavior. I wrote such a configuration and had the frontend output the potential at points increasing in distance from the center of the dipole, allowing me to fit the data to a $1/r^2$ curve.

2.7 Optimization and Multithreading

The purpose of this program is to not only be correct, but to be correct quickly. In order to achieve this, I have tried three things in combination:

1. General program optimization, targeted towards 64-bit x86 machines
2. The use of the x86 SIMD instruction set
3. The use of multithreading to split up the work among CPU cores

The first of these is pretty basic. The main part of the solver is written in C so that it can run directly on the processor. When the program is compiled, I instruct the compiler to optimize the program to the best of its ability, and the program is written in such a way that it can be optimized well and does not waste time doing unnecessary work.

As an example of this point, consider a 2-dimensional grid (2-D array), which must be iterated over (much like this program does). There are two ways to do this: either iterate row-by-row, or column-by-column. The end result is the same, and one might think the choice is arbitrary. However, when I tested the two different approaches, I found the row-by-row approach to be around 5 times faster. This is because of caching effects and the way that the CPU fetches data from RAM[7]. For a more detailed description of this, see Appendix B.

Another example involves data structure organization. Say a program needs to store an array of objects, and each object has data associated with it (for example, current value, previous value, error, initial state, etc). One might be tempted to write something like in Listing 1. The problem here is that if each object's `CUR_VAL` are all calculated together (as is often the case, and is definitely the case for this solving program), then the location of

the values that need to be loaded in succession by the processor are far apart. This, again, comes down to caching effects which are further explained in Appendix B.

```
1 struct {  
    double cur_val, prev_val;  
3    float err;  
    int initial_state;  
5 } objects[N];
```

Listing 1: Array of structures organization.

Fortunately, this also has an easy to make change that drastically improves performance: change to structure of arrays organization, thus reorganizing the memory layout so that related and commonly accessed-together values appear adjacent in memory[6]. This is shown in listing 2. While it may be a less intuitive way to describe and program a simulation, the result is significantly faster code the majority of the time.

```
1 struct {  
    double cur_val[N];  
3    double prev_val[N];  
    float err[N];  
5    int initial_state[N];  
} objects;
```

Listing 2: Structure of arrays organization.

2.8 SIMD Instructions

I have also made use of the SIMD (Single Instruction Multiple Data) instruction set. These instructions are available on modern 64-bit x86 processors^[5]². Their purpose is to enable what's known as vector processing of data on x86 machines. This provides the ability to apply a mathematical operation to one or more arrays, and have that operation be applied to each element. For example, I can have a 128-bit register filled with 4 values, $x = (1, 2, 3, 4)$, and another register like it, $y = (10, 10, 11, 11)$. I can then add them together to get $(11, 12, 14, 15)$, and this addition is issued with a single instruction. Furthermore, this takes the same amount of time as a single addition, because the processor can do the 4 separate addition operations in parallel.

I have written in the use of SIMD instructions in the solver program in order to both reduce branches and to parallelize the calculations required when calculating the next iteration of a given cell. I provided the ability to disable this functionality and instead do a non-SIMD style calculation in order to get benchmarks. A requirement of the SIMD code was that it give exactly the same result as the non-SIMD version (otherwise the implementation of the math would be incorrect). For this reason, when talking about the results of a simulation, it does not matter if SIMD was used or not; the use of SIMD only affects performance.

An important note is that when fully optimizing code for a machine which it knows about, `gcc` may choose to output SIMD instructions in order to do computation even if the user does not do so explicitly. For this reason, in places where I refer to the mode of the simulation as being “non-SIMD”, what I mean is that it does not use my manual SIMD instructions but it may still use compiler-generated SIMD instructions.

²More instructions are added in most new releases of CPUs. There are different SIMD instruction sets available on different processor models. I have used the `sse` through `avx` instruction sets in this code.

2.9 Multithreading

Finally, I have added multithreading support to the solver. This is done by spitting up the grid into subregions which are contiguous in memory, and spawning N threads via the pthreads API. Each thread is then given a region to work on, and they coordinate through shared state in order to decide when the solution has been reached. The initial implementation showed extremely poor results with multithreading (often being drastically slowed than the single threaded version) due to the use of atomic shared variables. It was then re-written such that the code was unsafe but much faster. This is because each thread modifies a shared non-atomic variable that keeps track of the current total change per iteration after that thread itself completes a number of iterations over its region. This kind of shared memory access is much faster than coordinating the threads, but it is also technically undefined behavior in C. Fortunately, on x86, it is actually safe, however there is potential information loss because a thread could overwrite another thread if they are trying to update the value at the same time. A more detailed discussion of the code and why it is unsafe can be found in Appendix B.

2.10 Benchmarking Performance

Benchmarks were done by placing two calls to `gettimeofday(3)` around the core of the solver. This function has a granularity of 1 microsecond[4]. This code returns both the number of iterations done and the difference between the two `gettimeofday` calls, which are then used to calculate the iterations per second. In order to take into account error from unpredictable scheduling effect from the operating system, each benchmark was run 300 times, and the standard deviation and mean were calculated, along with 95% confidence intervals. Several different computers were used: computer A was an Intel Core i5-3570K with 4 cores running at 3.4GHz with 16GB of RAM. Computer B was an Intel Xeon E5620

with 4 cores with 2 threads each, running at 2.4GHz, with 24GB of RAM. Computer B did not support all of the SIMD instructions that I used, so it was only able to benchmark the non-SIMD version of the code.

3 Results

All results describing the correctness of the simulation or discussing the closeness of the simulation to an analytical result are done using the single-threaded mode. This is because this makes the result of the simulation deterministic, and so it does not matter which machine the test was run on. For results that measure performance, the machine and options for the simulation used will be specified. There are times when the number of threads affected the correctness of the result.

Additionally, all results presented use the successive over-relaxation method for calculating the solution. This is because it converges more quickly on a solution than Jacobi iteration. Both methods produce the same results (as long as the Jacobi iteration runs for long enough), so the correctness results are unaffected by the choice of algorithm. Secondly, the performance measurements measure iterations per second, which is also largely independent of the choice of algorithm. Since successive over-relaxation is so much more effective at quickly converging upon a solution, it is presented as the program's primary mode of operation, with Jacobi iteration provided as a secondary and deprecated option.

3.1 The $\sin(x)$ Potential

Figure 2 shows the output of the solver program in single-threaded mode when given the $\sin(x)$ along one side potential, plotted with the Python library matplotlib. The configuration has a grid size of 300 by 300. It matches very well with the analytical result, which is

$$\frac{\sinh(y\pi/n)}{\sinh(\pi)} \sin(x\pi/n),$$

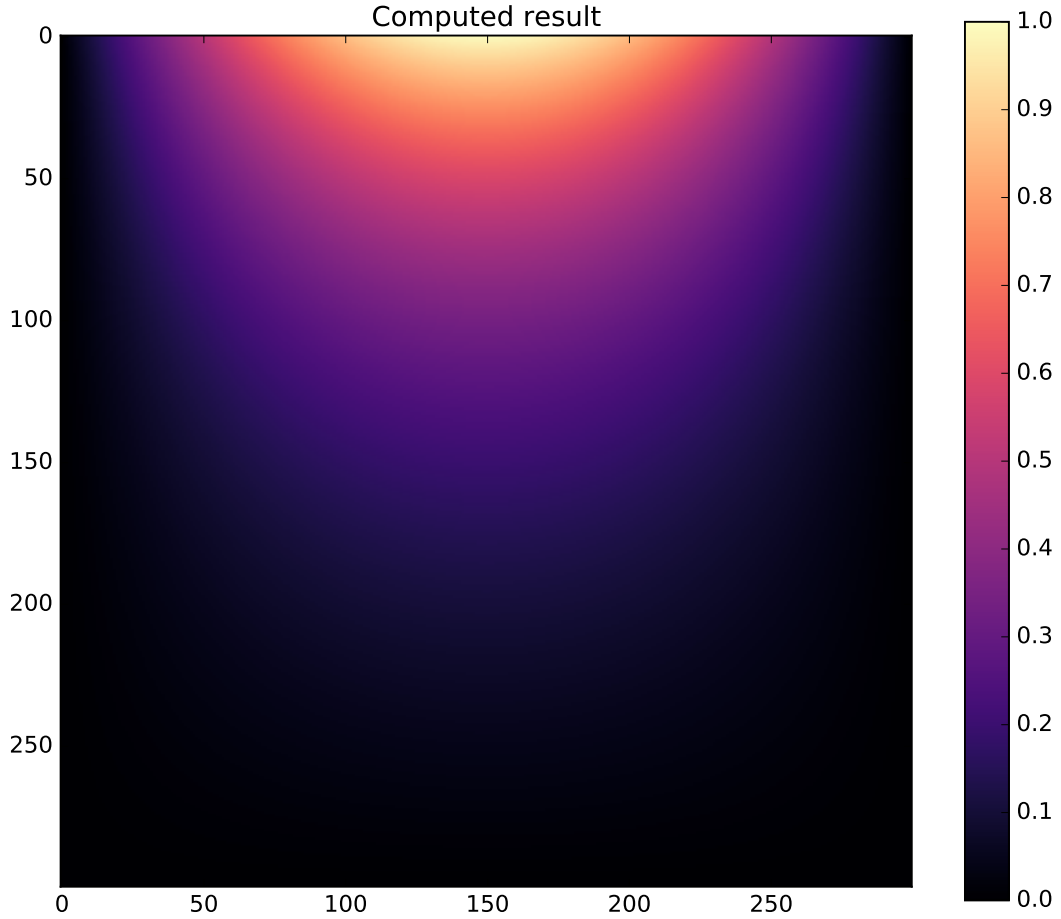


Figure 2: Result of simulation for the $\sin(x)$ potential.

and is shown plotted in a similar way in figure 3. The difference between the calculated result and the analytical result is shown in figure 4. The values shown in the difference map are small compared to the values in the result in all locations except the corners at the top, where the simulation does diverge somewhat, but never reaches values which are hugely divergence from the analytical result. Using the root-mean-square error statistic described earlier, we have calculated that the RMS error in this simulation to be SOME NUMBER.

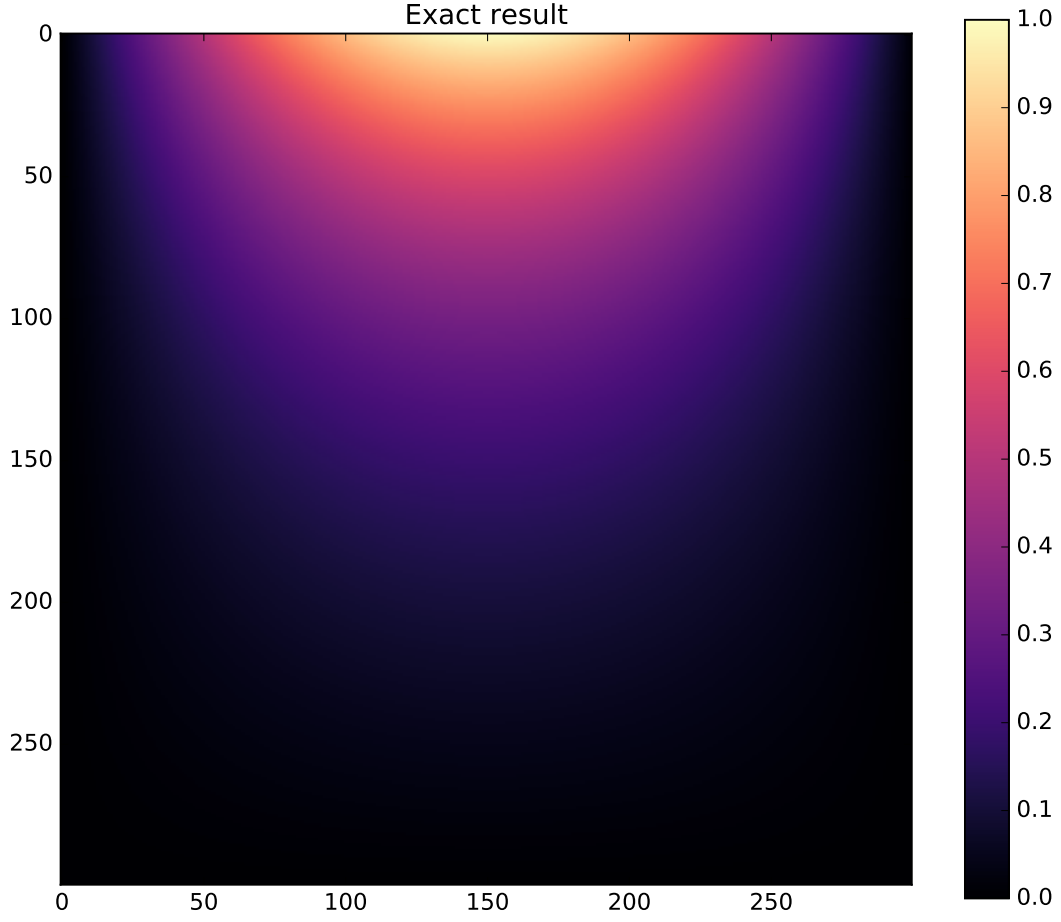


Figure 3: Analytical result for the $\sin(x)$ potential.

3.2 Fitting to a $1/r^2$ Dipole Potential

The dipole configuration had a grid size of 400 by 400, with all walls given a Dirichlet boundary condition of zero. Two charges were placed near the center 10 cells apart with an arbitrary but equal and opposite charge. The result of this simulation is shown in figure 5, with the addition of constant-potential contour lines. The simulation shows two opposite charges which near perfectly mirror each other. The vertical line in the center represents a

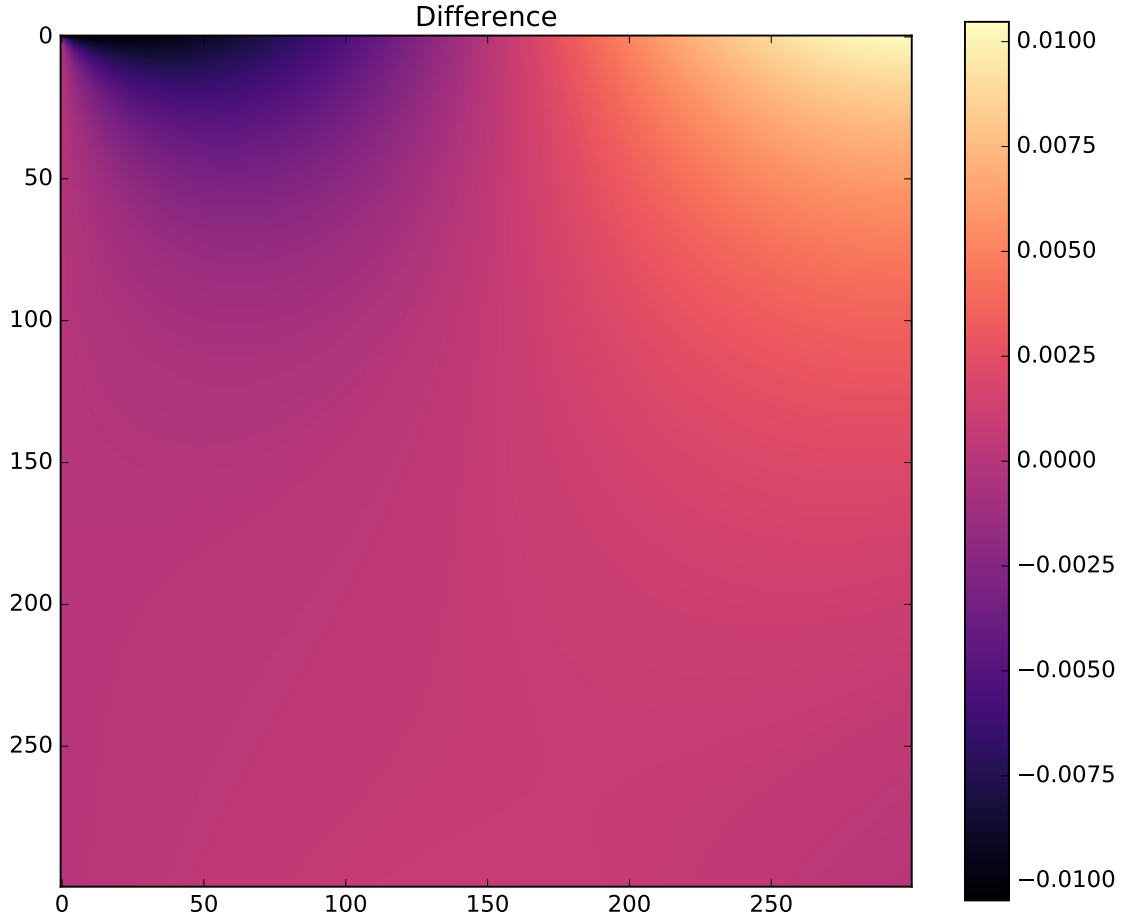


Figure 4: Difference between the calculated and analytical results for the $\sin(x)$ potential.

constant potential contour with a value of zero, which is expected from an electric dipole. This is another excellent verifier that the simulation is producing correct results.

The potential was then measured at locations increasing in distance from the dipole in an arbitrary direction. These data were then plotted using gnuplot, and a line of the form $y(x) = A + B/(r - r_0)^2$ was fit to the data, which is shown in figure 6 The fit is excellent, confirming that the program had properly simulated an electric dipole. The error bars on

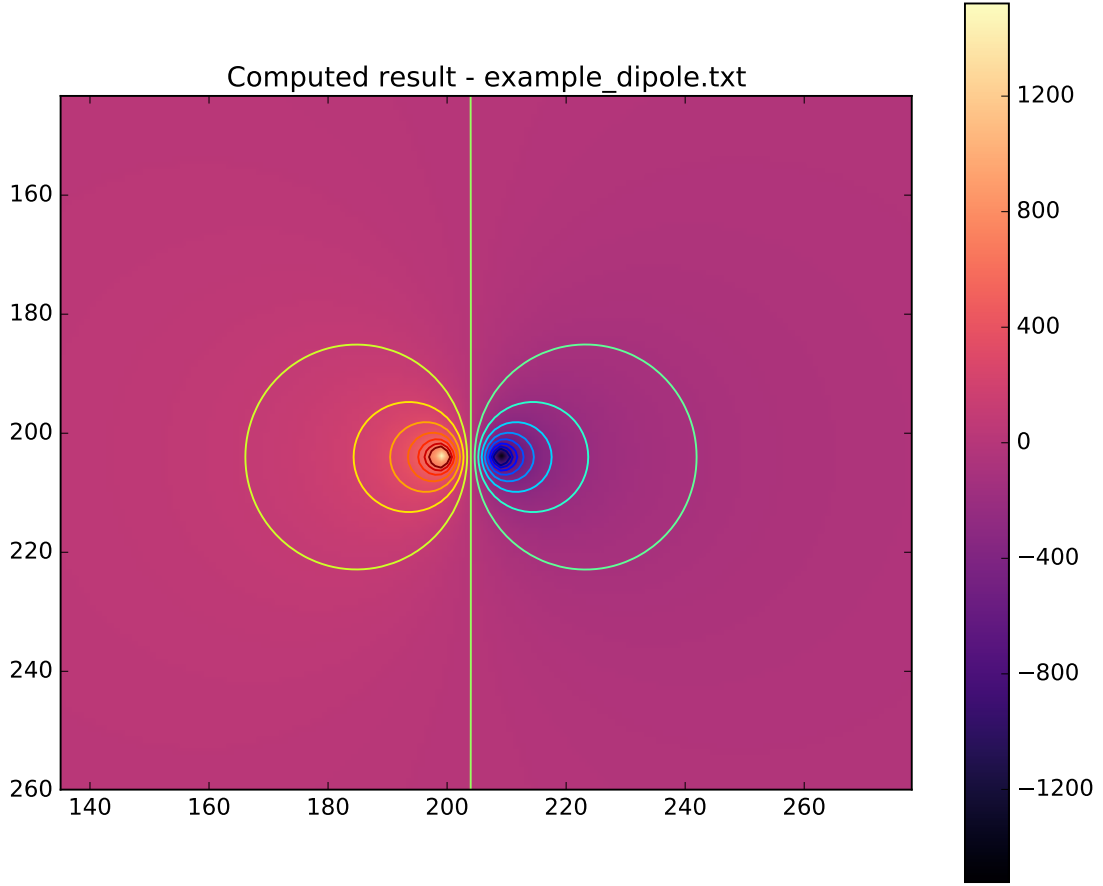


Figure 5: Result of the simulation for the dipole configuration, including contour lines.

the data points come from the difference between successive iterations at the end of the simulation of the dipole. They are very small, but non-zero.

The dipole is an example which lends itself to presenting another feature of the solver program. If passed the `-V` option, it will take the negative gradient of the result and plot that as a vector field. This has the effect of plotting the electric field of the result of the simulation, which for the dipole is shown in figure 7.

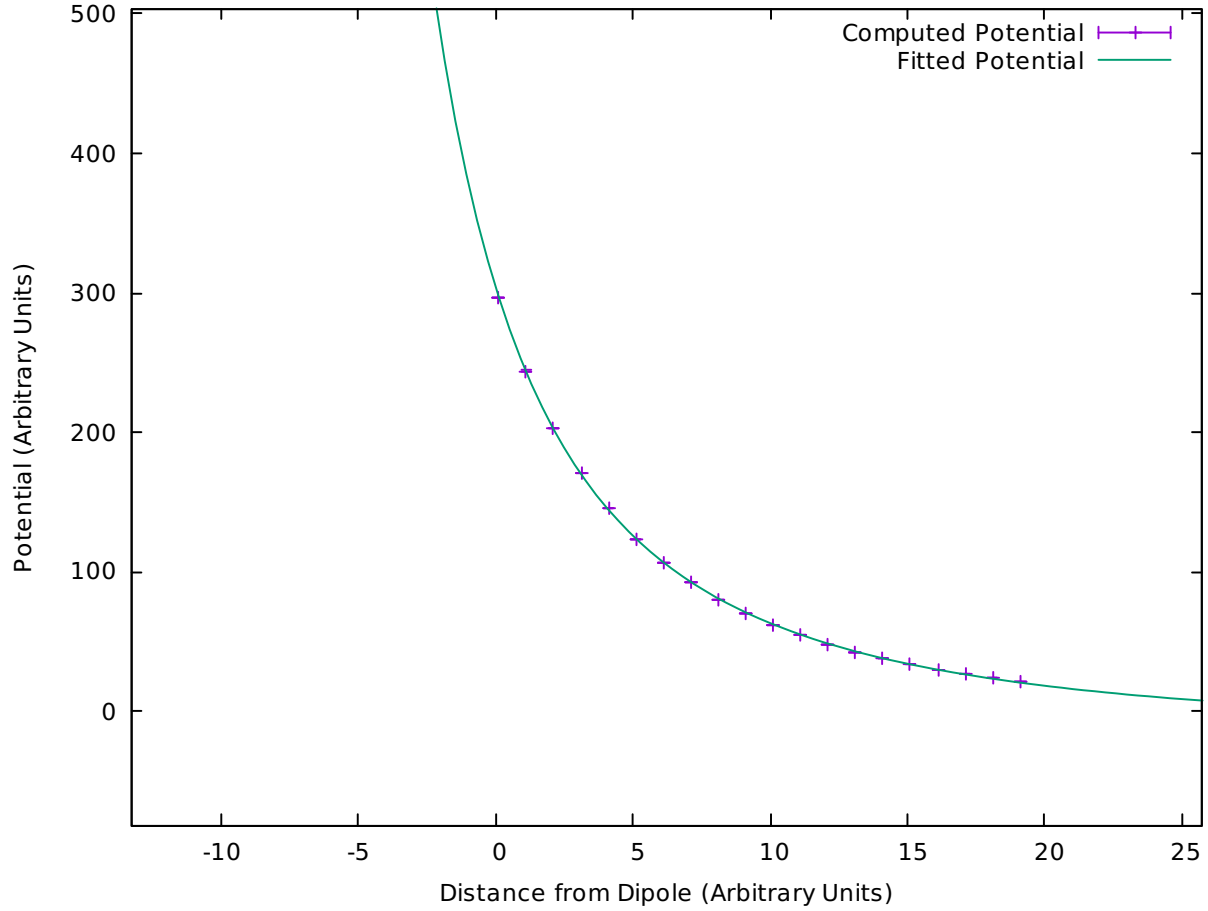


Figure 6: Fitting a $1/r^2$ curve to the calculated dipole potential.

3.3 Performance

Figure 8 shows the performance in iterations per second for a varying grid size for given configuration, in this case the $\sin(x)$ along a wall potential. The performance tests were run square grids of size 50 by 50, 300 by 300, and 1000 by 1000. For each grid size, 4 different simulation options were given: single-threaded non-SIMD, single-threaded SIMD, multi-threaded with 4 threads non-SIMD, and multi-threaded with 4 threads and SIMD.

In the case of the 50 by 50 grid, the performance was similar for all cases except for multi-threaded with SIMD. I believe the reason that the SIMD multi-threaded version was slower was due to the manual SIMD code being less cache-friendly, resulting in the threaded versions

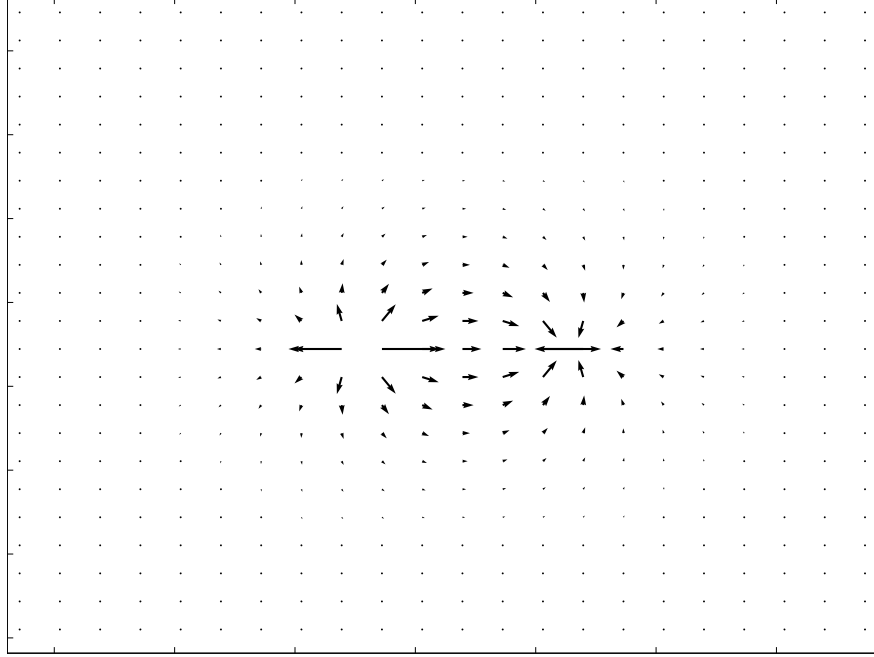


Figure 7: Result of program outputting the electric field for the dipole configuration.

having more contention. Both the 300 by 300 case and 1000 by 1000 case display similar performance behavior: the SIMD case is faster than the non-SIMD case when single threaded, and the SIMD case is slower than the non-SIMD case when multi-threaded. This can be explained in the same way as it was for the 50 by 50 case. Another characteristic is that the multi-threaded case (with 4 threads) is faster than the single threaded case. In fact, it is better than I expected it to be.

Additionally, the performance depends heavily on the grid size. This makes sense, as the amount of work needed to be done per iteration is proportional to the number of cells in the grid, and the number of cells in the grid is n^2 . Another note is that the variance on the performance is significantly higher for all of the multi-threaded cases. This is because having multiple threads increases the program's dependence on the scheduling quirks of the operating system. Finally, the variance for the 50 by 50 case multi-threaded versions are significantly higher than all other variances for the other multi-threaded cases. I believe that

this is because the overhead of managing threads compared to how much work they actually have to do is large in this case compared to the other cases.

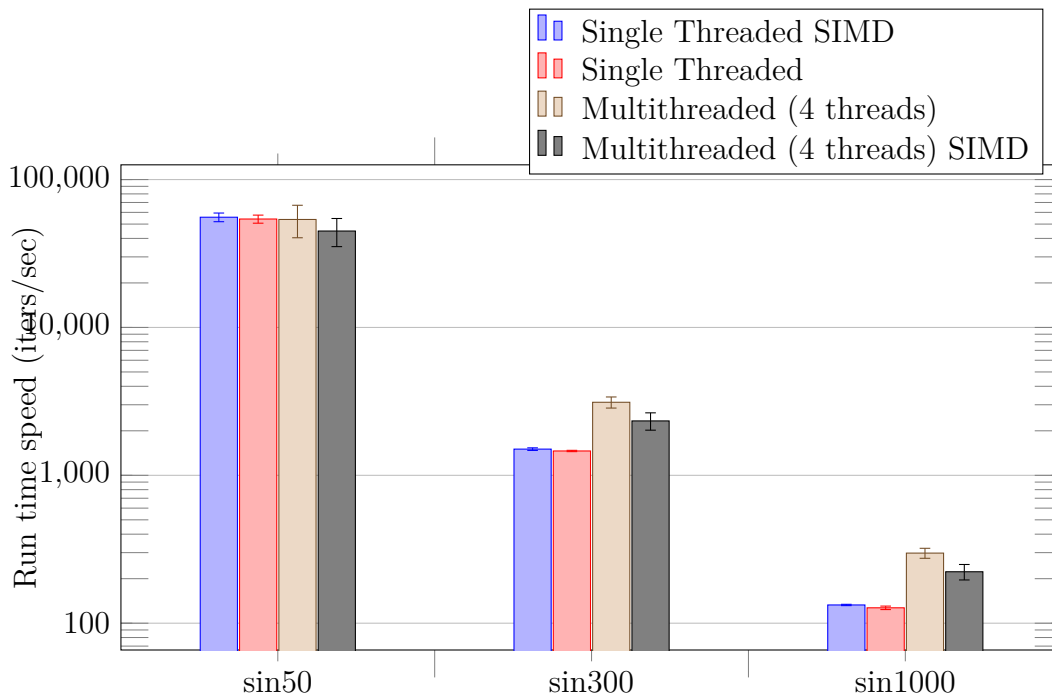


Figure 8: Different grid sizes for the $\sin(x)$ configuration, and their performance characteristics.

Figure 9 shows the performance in iterations per second of the solver when run with the $\sin(x)$ potential with a grid size of 300 by 300. This was done on computer B, and so was done only in non-SIMD mode. The results show a significant improvement in performance when using multiple threads compared to the single threaded case. One may naïvely have expected the performance to scale linearly with the number of threads; however, this would be very unlikely. As the number of threads goes up, the amount of work that can be done in parallel would seem to scale with the number of threads, until one considers processor cache effects (specifically cache coherency and cache size limitations). This is seen in the case of 4 threads, as it is not 4 times as fast as the single threaded case, only about 3.2 times as fast. Increasing the number of threads quickly starts giving diminishing returns, as the 8 threads case is now only slightly better than the 4 threads case. Recall that computer B, on which these tests are run, has 4 cores with 2 threads each, totally 8 possible threads running

in parallel. After the 8 threads case, we start to see a drop in performance; the 16 and 32 threads case are worse than the 8 threads case. This is because the machine cannot run more than 8 threads in parallel, so in order to have 16 threads it must rely on the scheduling of the operating system in order to get work done. In a case of many threads, each doing heavy computation and memory accesses (as is the case here), it does not usually help to increase the number of threads beyond what the machine can handle, as is shown here.

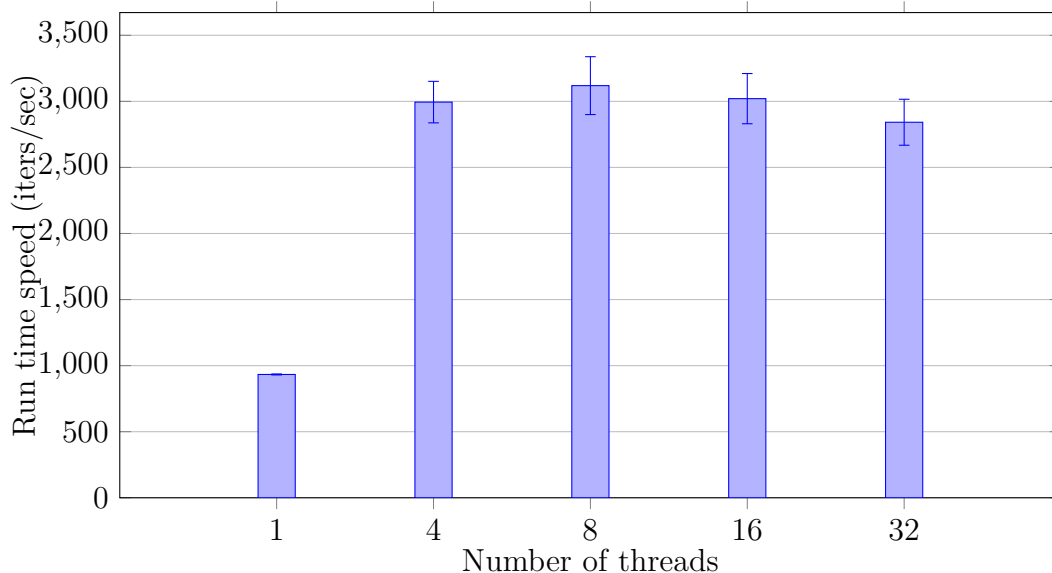


Figure 9: Performance results for a varying number of threads in non-SIMD mode. These tests were done on computer B, using the $\sin(x)$ potential with a grid size of 300 by 300.

Furthermore, figure 10 shows that it is disadvantageous to run the simulation with a large number of threads. The RMS error of the simulation compared to the analytical result increases rapidly with the number of threads. This is also reflected in the result of the simulation, which is noticeably wrong in the case of 16 or higher threads. I believe this to be the result of the way the threads are synchronized, or more accurately, how they are not. The threads run mostly in isolation, sharing the grid and doing operations on it without waiting for other threads to complete any of their work. This means that if one threads runs faster than the others, it may complete a different number of iterations than the other threads in a given amount of time. This would result in some sections of the grid receiving more iterations than others, which could have the effect of invalidating the simulation. The

alternative method would be to synchronize the threads, and have them wait for each thread to complete a given iteration before moving on to the next one. Indeed, this was the original design of the system, and I found it to be so slow compared to single threaded that I changed the design around to the less accurate but much faster design that is shown here.

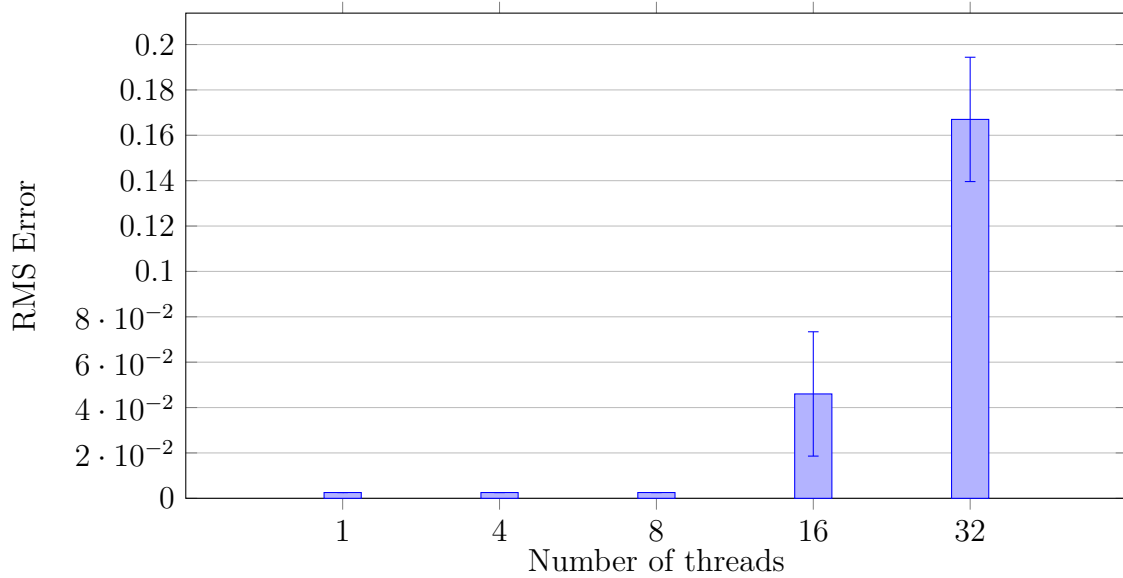


Figure 10: Error results for a varying number of threads in non-SIMD mode. These tests were done on computer B, using the $\sin(x)$ potential with a grid size of 300 by 300.

4 Conclusion

This program correctly solves certain boundary value problems, enabling simulation of complex electrostatics situations. It does this in an efficient manner by leveraging proper programming languages and techniques and utilizing the basics of modern processor design. There were numerous surprises found in developing this software and looking at the resulting physics that it described.

4.1 Correctness

The most important aspect of this program is that it produces an accurate result, no matter how fast it runs. I was able to use a problem for which I had an analytical solution to confirm that the simulation was producing accurate results and found that the results were of high quality for the single threaded and low number of threads cases. Additionally, I was able to use the solver program to accurately calculate that the electric potential from an electric dipole drops off as $1/r^2$, further confirming the correctness of the program.

4.2 Performance

Since the simulation is correct, I was able to then turn my attention towards improving the performance without reducing the accuracy of the simulation. Most of this was easy to do without reducing accuracy, as it was simply transforming one valid implementation of the code into another and measuring to ensure that my new implementation did in fact cause an improvement. Most of this was done by improving cache locality and reducing memory accesses. I have left out a lot of this process for brevity, however an overview of it and the concepts used can be found in appendix B.

The next significant performance gain attempt that I had made was the use of SIMD instructions. I was hoping for a much more significant performance boost than I saw, which was rather miniscule in reality. The only small improvement was not surprising in the single threaded case because I knew that `gcc` utilizes SIMD when it can, but I did not expect to see the manual SIMD instructions degrade the performance in the multi-threaded case. I speculate that this is caused by my manual SIMD code being less cache friendly, which is okay in the single threaded case and allowed a small performance improvement, but in the multithreaded case the threads are all contending for limited cache space, so the additional memory usage slowed them down.

The multithreading actually improved the performance significantly more than I expected it to. I had predicted that the algorithm would be cache-bound due to the large amount of memory accesses that it makes. This often results in insignificant performance improvements (or downright performance degradation) when adding multiple threads to a single threaded program. Of course, the cardinal rule of any scientific experiment (including attempting to optimize a program) is to measure and compare results. Clearly the multi-threaded case is significantly better in performance, with a trade-off in accuracy, which may certainly be acceptable.

4.3 Future Work

There is still a significant amount of optimization work that can be done on this program. A careful reading of Intel optimization manuals would likely provide numerous ideas for how to further improve the code for the program. The manual SIMD code could also be improved, either by utilizing more up-to-date instruction sets or by improving the existing algorithm. The multithreaded code could also be improved, by allowing the threads to communicate in a limited way in order to better synchronize their work to improve the quality of the

result without hopefully degrading the performance too drastically. An additional approach could be to use the multithreaded case for a few hundred iterations before switching over to the single threaded version in order to smooth out and cleanup the results. This would have a significant advantage for large simulations, as the multithreaded code could get the result most of the way before letting the single threaded code produce an accurate result in a shorter amount of time than if the single threaded code ran from the start.

Another significant way in which to improve the real world performance of the program is in addition to optimizing the speed at which the program can calculate an iteration, to reduce the number of iterations required to converge on a solution. Indeed, this is what the use of successive over-relaxation achieves, but there may be ways to further improve this. One such method could be to improve the initial “guess” of zero everywhere by dividing the grid into a much coarser grid and running a few iterations on that before returning to the full grid and completing the simulation as before.

Appendices

A Program Usage

Compiling the engine (the C library) requires `make` and `gcc` (or `clang`, although `clang` is untested), and can be done by simply invoking `make`. This will also create an executable file, `solve`. This is a script that invokes the python example frontend so that it may be used.

For more details on using the python frontend or the C backend engine, see the man pages (reproduced on the following pages). There are several example files, including `example_sin.txt` and `example_dipole.txt`. For more details on writing configurations for a desired simulation or experiment, see the man pages.

The source code also contains a directory named `tests`, which contains python code to generate verifier data files for the `example_sin` configuration. Finally, all of the data written on in this document are present in the source directory, under `results`. The source code for the C library is in `engine`, and the source code for the python example frontend is under `frontend`. The source code and git revision history is available at <http://github.com/dbittman/statics-solver>.

B Computer Program Optimization

When optimizing a program, there is an important cardinal rule to follow which at times is easy to forget in the face of increasingly interesting and complex minor optimizations designed to squeeze every bit of performance out of your computer. Unfortunately, due to the complexity of modern computer hardware, it is very difficult to predict if a given micro-optimization will actually be beneficial or not. For this reason, the most important rule when taking a working program and making it fast is³ *benchmark everything*.

The field of program optimization is huge. There is a small number of people who, for a given platform, are really *good* at optimization. There is an even smaller number who are truly experts. I do not fit into either of these categories. For this reason, the program that I have written here is *decently* optimized. I believe that the maximum throughput is not orders of magnitude better than what I have achieved, however, there is still a lot of work that can be done (including a potentially significant optimization that I have neglected, and will talk about later).

B.1 Lazy Programs Are Better

Before taking up a sharp cutting tool and carving away a program to make it run faster, it helps to examine what exactly is the program trying to accomplish and how. There are often multiple ways to solve a problem, and some of them may be more efficient than others. In my case, I utilized two different methods of solving a problem: Jacobi iteration and successive over-relaxation. One of them was significantly faster than the other not because each iteration of the program was quicker but because it had to do less iterations. It was using a smarter algorithm.

³besides *don't break it*, of course. An incorrect program is worthless, no matter how fast it is.

Computer scientists talk about algorithmic complexity using order notation. For example, given a random list of numbers, if I want to test to see if a certain number is in the list, I need to look through the whole list in order to check if it is there. This is an $O(n)$ operation, meaning that if the list is of size n , I have to do n operations to determine if the item is there or not. Data structures commonly have operations associated with them, and those operations (such as `find`, `insert`, etc) have associated time costs. Choosing the correct data structure for a problem is one of the first steps towards correctly organizing a program for it to be efficient. If a program is constantly searching through a list to find values when instead a simple hash table could be used, the program will run slowly. A program should not do additional work when it does not need to. Only once the data structures have been properly designed and the algorithms properly thought through should *optimization* be done.

B.2 Freebies

There are some optimizations which can be used in high confidence and are also very easy to do. The first, and most obvious, is to use a fast language. While writing in `C` may be a chore compared to writing in `python`, the resulting code will be much faster. This is the reason behind the organization of this program: the frontend (which has no performance requirements) is written in `python` because parsing the configuration file is much easier. The backend needs to be fast, so it is implemented in `C`.

A quick and rough experiment can easily convince you that this is true (remember how I said to benchmark everything?). Credit to my friend Chris Milke for doing this test. Write the following code in both `python` and in `C`: initialize a variable to zero. Then loop from zero to some large number (we used 123,456,789), and add that iteration to the variable. After the loop, print the variable. The resultant `python` program took 14.25 seconds to run on my computer, but the `C` code completed *instantaneously*. Why is `python` slower than `C`

in general? Because python is an interpreted language. The code that you write is read by another program and executed by that program. In contrast, C code must be compiled into machine code. This is then run directly on the processor, which cuts out the middle-man of the interpreter.

When using C, there are some more things that you can get for free: compiler optimizations. When compiling a C program, the command looks something like this: `gcc -Wall -Wextra foo.c`⁴. This will result in a compiled C program (a binary file, containing the machine code), but it will be unoptimized. There are reasons why one may want an unoptimized binary (they are typically easier to debug, for example), but generally when you compile your program for actual usage you will want to optimize it. An aggressive optimization compilation command may look something like: `gcc -Wall -Wextra -O3 -ffast-math -mavx`, as a start. The main flag here is `-O3`, which enables almost every optimization that gcc can do.

Going back to the example from earlier, the optimized version of the C program completed instantly, but the unoptimized version took 0.33 seconds. The unoptimized version is still 43 times faster than the python code, but what was the optimizing compiler doing to make it so much faster with `-O3` on? Part of optimizing is understanding why something is faster, so lets look at the assembly code. On UNIX, this can be done with the `objdump` command. Table 1 shows the annotated assembly from the function `main`. If you don't know x86_64 assembly language, you will at least notice that the optimized version is much shorter. If you do know how to read the assembly, you'll notice that the unoptimized version is doing almost literally what the C code says to do: set a variable to zero, and iterate from zero to a big number, adding that iteration to the variable, and finally printing the variable. The optimized code just prints a large number – which happens to be the result of the sum.

⁴as an aside, you should **always** compile with `-Wall` and `-Wextra`, and eliminate all warnings from your program. Warnings are warnings for a reason, and should rarely be ignored.

Unoptimized	Optimized
<pre> mov [rbp-0x10], 0x0 jmp .loop_end .loop_top: mov rax, [rbp-0x10] add [rbp-0x8], rax add [rbp-0x10], 0x1 .loop_end: cmp [rbp-0x10], 0x75bcd14 jle loop_top mov rsi, [rbp-0x8] mov edi, 0x4005b4 mov eax, 0 call 4003c0 <printf@plt> </pre>	<pre> movabs rsi, 0x1b131147ee6b52 mov edi, 0x400594 xor eax, eax jmp 4003c0 <printf@plt> </pre>

Table 1: Unoptimized and optimized assembly from the sum-a-lot-of-numbers example.

This is just a simple example, but a good C compiler can drastically improve the performance of a program. Here, it made a program that took a third of a second (which is already drastically faster than the interpreted python version!) finish instantly.

B.3 A Model of Processor Design

When first learning to program in a language that does more accurately display what is going on inside of a computer, it is helpful to have a mental model for what is going on during each simple operation that a programmer might want to do. Since programs are sets of instructions which define transformations on data, it is tempting to think about a model such as in figure 11a, where there is some nebulous device which follows your instructions and is connected to a large amount of memory in which your variables and data structures live.

Unfortunately, while this is a convenient abstraction to think about programming up an application, it is not actually what is going on⁵. A more realistic model is shown in figure 11b. There are several key take aways from this. One of them is that RAM is *really slow*. The second is that the processor chip is aware that RAM is really slow and so it aggressively caches everything that it ever gets out of RAM. This diagram shows the actual detailed level of caches that cores have on them, and that they have a shared cache, but most of that complexity is not needed for a basic understanding. When one makes first steps towards optimizing a program, it is vital to think in the more complex model, because a huge amount of time will be spent trying to optimize the program according to how the processor actually works.

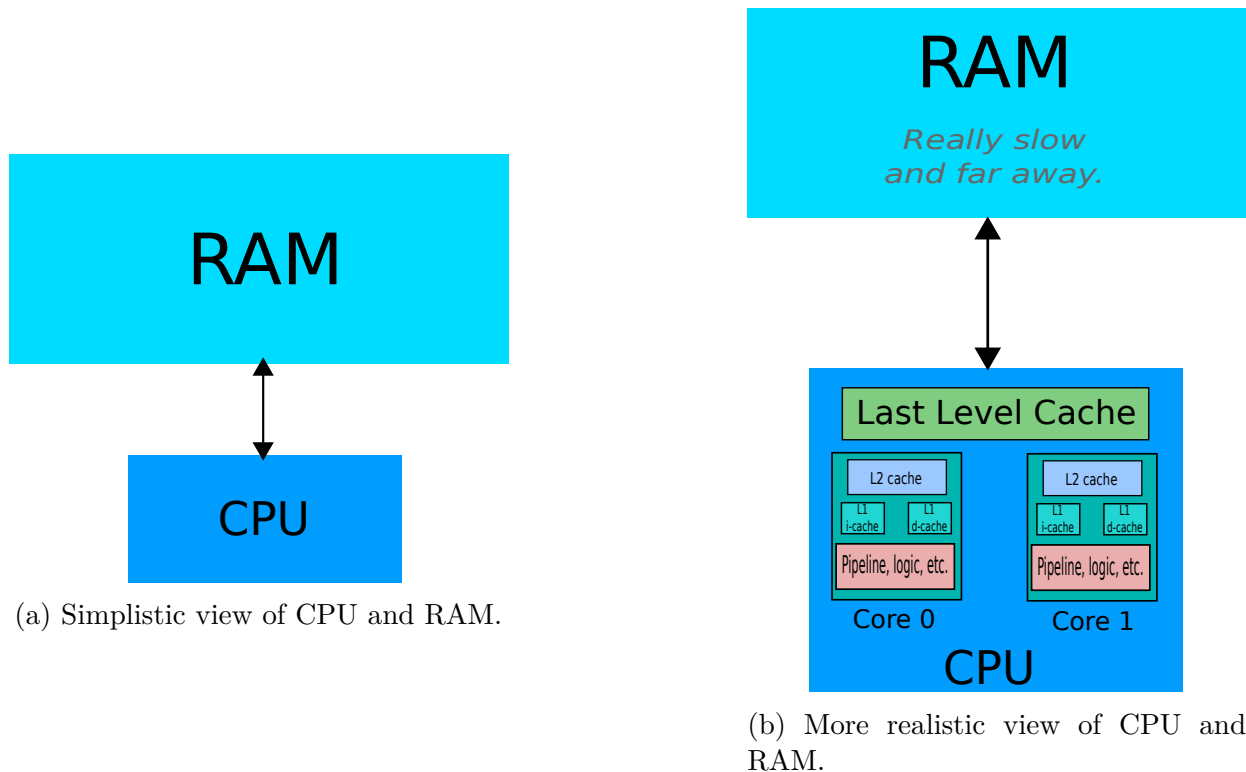


Figure 11: View of programming that programmers like to have (a), versus the the more realistic and complicated view (b).

⁵I do not claim that what I am about to describe is what is “actually going on” either, but it is at least significantly less of a lie.

B.3.1 Caching

On modern processor chips, the vast majority of the silicon is used to implement a memory cache. This gives an indication of how important it is. Think of the cache as a smaller, but much much faster memory that the processor accesses whenever it would access RAM. If it finds the value in the cache, reading a value can take between 1 and 10's of nanoseconds. However, if it does not find the value in the cache, it must then perform a read from memory, which can take hundreds of nanoseconds. For scale, a single instruction very often takes less than a nanosecond to do. This means that if a program makes a lot of memory accesses, the processor will spend a lot of time waiting and not much time calculating.

Earlier in this document, I described a phenomenon which describes how important caching is. Given a 2-D array, there are two ways to iterate over the whole thing: line-by-line or column-by-column. On my computer, the line-by-line method was five times faster than the alternative. This has to do with the cache, and specifically, *cache lines*. When reading a value out of memory for the first time, the processor actually reads many values, something like 64 bytes worth of data. This is fact of the design of the hardware, and is partly because if you read a value x from RAM, the processor expects you will want to access memory that is nearby x in the near future. Thus, if you iterate over an array line-by-line, when you access the first element of a line you are actually reading in multiple elements of the array into the cache at once. When the processor then tries to access the second element of the line, it finds that it is already in the cache, making that access much faster than if it had not been in the cache. If you iterate over the array column-by-column, you read the first element of the line, which loads several elements from that line into the cache, and then *those values are ignored* and the first element of the second line is read. This means that when reading the first element of the lines, each time the processor has to fetch the value from RAM and does not get a chance to make use of the cache. By the time an entire column has been read, the cache may be full, forcing the processor to remove old values from it. Then, when the

program reads the second element of the first line, that data no longer in the cache. This is extremely slow.

If possible, try to keep the amount of memory that the program accesses small. Try to organize the access patterns such that it is something that the processor is good at handling (like the iteration direction choice above). I have utilized these methods in my solver program. An example for keeping memory footprint small was to reuse variables inside the grid that I knew were safe to reuse. Additionally, I used `floats` instead of `doubles` wherever possible, as this reduces memory usage as well.

B.4 Structure of Arrays

B.5 Multithreading

B.6 Vectorization

B.6.1 Branch Reduction

C Program Implementation

References

- [1] Mary L Boas. *Mathematical methods in the physical sciences*, volume 2. Wiley New York, 1966.
- [2] Richard J. Gonsalves. Poisson’s equation and relaxation methods, 2011.
- [3] David Jeffrey Griffiths and Reed College. *Introduction to electrodynamics*, volume 3. prentice Hall Upper Saddle River, NJ, 1999.
- [4] IEEE / The Open Group. *gettimeofday(3) Linux Users Manual*, 2.23 edition, May 2016.
- [5] Intel Intel. and ia-32 architectures software developer’s manual. *Volume 3A: System Programming Guide, Part, 1*, 64.
- [6] R Intel. and ia-32 architectures optimization reference manual. *Intel Corporation, May*, 2007.
- [7] R Intel. and ia-32 architectures software developer’s manual volume 1: Basic architecture. *Intel Corporation*, 64.
- [8] Aleka McAdams, Eftychios Sifakis, and Joseph Teran. A parallel multigrid poisson solver for fluids simulation on large grids. In *Proceedings of the 2010 ACM SIGGRAPH/Eurographics Symposium on Computer Animation*, pages 65–74. Eurographics Association, 2010.
- [9] Marc Spiegelman. Myths and methods in modeling. *Columbia University Course Lecture Notes*, available online at <http://www.ldeo.columbia.edu/~mspieg/mmm/course.pdf>, accessed May 2016, 6:2006, 2004.